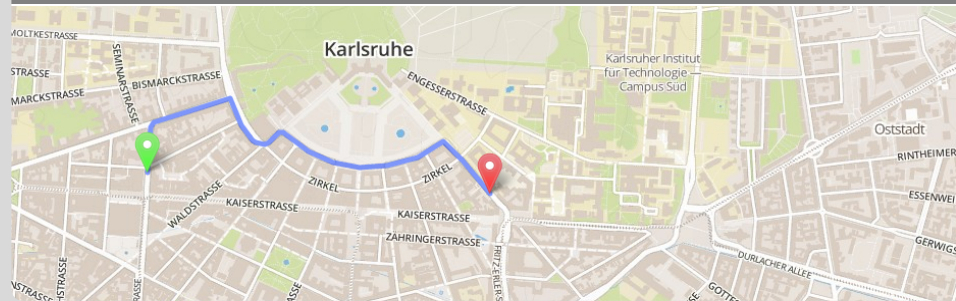


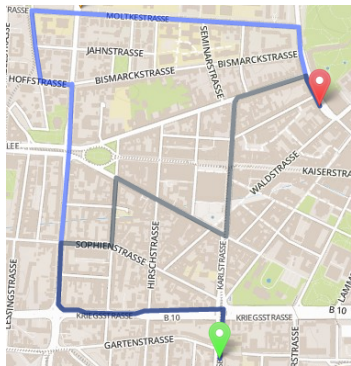
Engineering Parallel Bi-Criteria Shortest Path Search

Stephan Erb – stephan.erb@student.kit.edu

Institute of Theoretical Informatics




The Problem



Simultaneous optimization of
multiple (possibly) conflicting
objectives
(e.g., **travel time** and **highway tolls**)

Bi-Criteria Shortest Path Problem




Graph $G = (V, E)$ with 2-dimensional edge weights $\vec{w} = (x, y)$

Pareto optimal path are non-dominated

A path p is dominated by a path q iff
 $(q \neq p) \wedge (q_x \leq p_x) \wedge (q_y \leq p_y)$

Bi-Criteria Shortest Path Problem




Graph $G = (V, E)$ with 2-dimensional edge weights $\vec{w} = (x, y)$

Pareto optimal paths are non-dominated

A path p is dominated by a path q iff
 $(q \neq p) \wedge (q_x \leq p_x) \wedge (q_y \leq p_y)$

Goal: Find all Pareto optimal paths from a start node to all other nodes
(one-to-all)

The Parallel Pareto Search Algorithm



Parallel Label-Setting Multi-Objective Shortest Path Search by P. Sanders and L. Mandow

IEEE International Parallel and Distributed Processing Symposium (IPDPS2013)



The Parallel Pareto Search Algorithm



Parallel Label-Setting Multi-Objective Shortest Path Search by P. Sanders and L. Mandow

IEEE International Parallel and Distributed Processing Symposium (IPDPS2013)



Provably efficient for two dimensions.
But is it also practical?

Dijkstra's Algorithm

Procedure DijkstraSearch(G, s)

$L[v] = \infty$ for all $v \in V$; $L[s] = 0$

PriorityQueue $Q = \{(s, 0)\}$

// tentative labels

while Q is not empty **do**

$(u, l) = Q.removeMinimum()$

foreach $(u, v) \in E$ **do**

$l' = l + weight((u, v))$

// candidate label

if l' represents new best path to v **then**


 update $L[v]$

 update Q


Label Setting Property:

Shortest path distance of node/labels extracted from Q is final

Bi-Criteria Label Setting



Bi-Criteria Label Setting



Classic Approaches [1]:


Extract **any** Pareto optimal label from Q

Sanders & Mandow:

Extract **all** Pareto optimal labels from Q

The Parallel Pareto Search Algorithm

All steps are
parallelized.



Goal:

Apply sequence of k **sorted updates** (insertions & deletions)

Proposed Solution:


Parallel **red-black tree** of Frias and Singler [2]

Problem:


Binary trees can be multiple times slower than
cache-sensitive B-trees [3]

Parallel Bulk Updates for B-trees

Updates:
(sorted)




B-tree:




Parallel Bulk Updates for B-trees

Updates:
(sorted)

B-tree:




Parallel Bulk Updates for B-trees



Subtrees updated **independently in parallel**.


Parallel Rebalancing of B-trees

Parallel **partial-rebuilding** [4] of unbalanced subtrees:




Amortized update bound of $\mathcal{O}(k/p \cdot \log N)$ for the application of k updates using p threads on a tree of size N .

Experimental Results



The Parallel Pareto Search Algorithm

All steps are
parallelized.



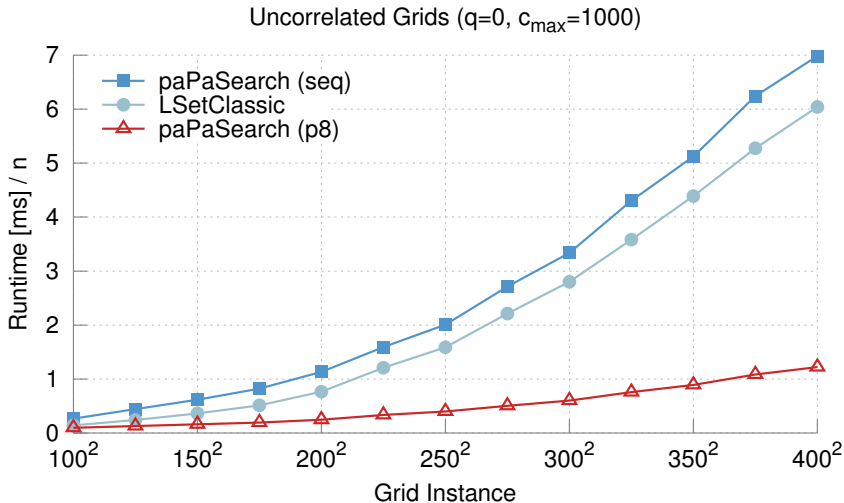
System:

- 2× Intel Xeon E5-2670 with **16 cores** in total, **2.6 GHz**,
2× 20MB L3 Cache
- 64 GB main memory

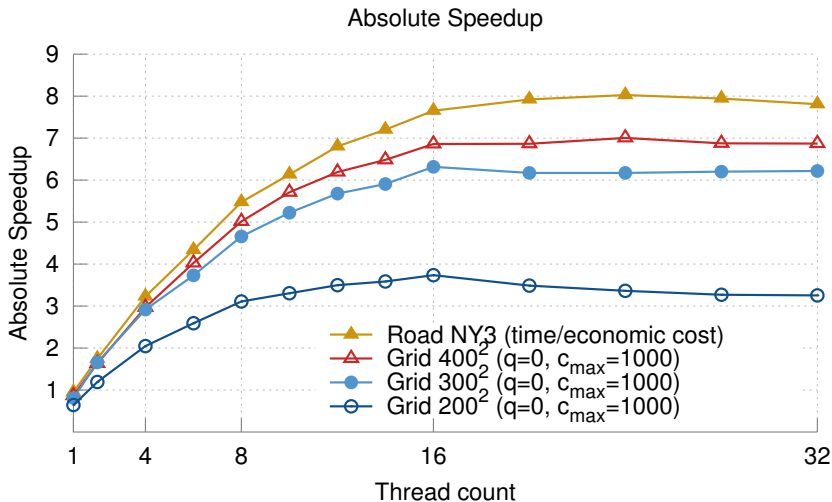
Tuned Competitor:

- highly tuned binary heap [5]
- fast dominance checks

Experimental Results




Experimental Results




Component Analysis

paPaSearch Components on Road NY₃ (time/economic cost)

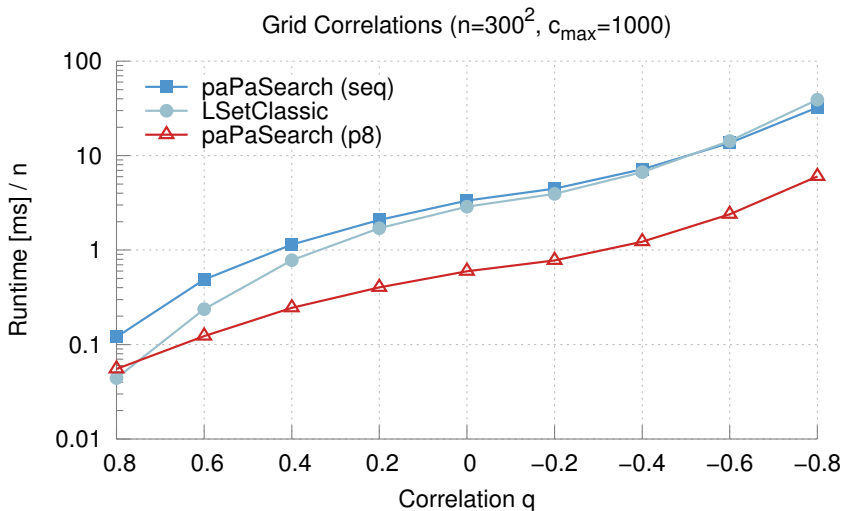


Component Analysis

paPaSearch Components Speedup Road NY₃ (time/economic cost)



Experimental Results



Summary:

- Parallel label-setting is **practical** on modern **multiprocessors**
- It excels at **large & difficult instances**

Future Work:

- Delta-stepping [6] to increase work per iteration
- scalable sorting of small collections?
- Point-to-point search

Thank You!


Results will be published on
<http://algo2.itl.kit.edu>.

- [1] P. Hansen, “Bicriterion path problems,” in *Multiple Criteria Decision Making Theory and Application*, ser. Lecture Notes in Economics and Mathematical Systems. Springer Berlin Heidelberg, 1980, vol. 177, pp. 109–127. [Online]. Available: http://dx.doi.org/10.1007/978-3-642-48782-8_9
- [2] L. Frias and J. Singler, “Parallelization of bulk operations for stl dictionaries,” in *Euro-Par 2007 Workshops: Parallel Processing*, ser. Lecture Notes in Computer Science, L. Bougé, M. Forsell, J. Träff, A. Streit, W. Ziegler, M. Alexander, and S. Childs, Eds. Springer Berlin Heidelberg, 2008, vol. 4854, pp. 49–58. [Online]. Available: http://dx.doi.org/10.1007/978-3-540-78474-6_8


- [3] T. M. Chilimbi, M. D. Hill, and J. R. Larus, “Cache-conscious structure layout,” in *Proceedings of the ACM SIGPLAN 1999 conference on Programming language design and implementation*, ser. PLDI '99. New York, NY, USA: ACM, 1999, pp. 1–12. [Online]. Available: <http://doi.acm.org/10.1145/301618.301633>
- [4] M. Overmars, *The design of dynamic data structures*. Springer Verlag, 1983.
- [5] P. Sanders, “Fast priority queues for cached memory,” *Journal of Experimental Algorithmics (JEA)*, vol. 5, Dec. 2000. [Online]. Available: <http://doi.acm.org/10.1145/351827.384249>
- [6] U. Meyer and P. Sanders, “ δ -stepping: a parallelizable shortest path algorithm,” *Journal of Algorithms*, vol. 49, no. 1, pp. 114 – 152, 2003. [Online]. Available: <http://www.sciencedirect.com/science/article/pii/S0196677403000762>

Backup

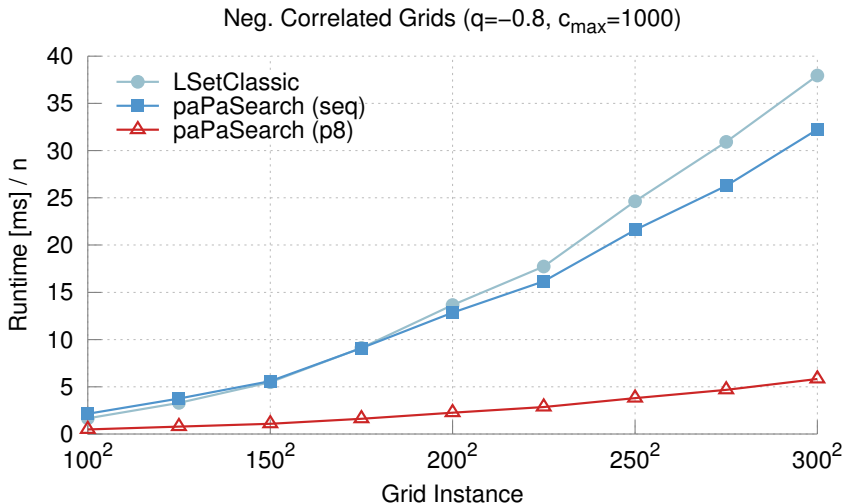
Experimental Results

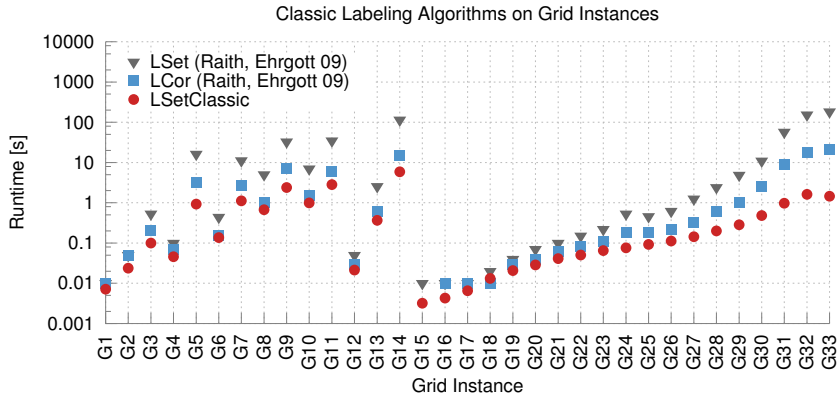



Experimental Results




Experimental Results







Fast dominance checks on lexicographically sorted label sets:




Predecessor queries as **tree traversals**.
Running in $\mathcal{O}(\log N)$ for a label set of size N .

Tuned Dominance checks

Batched dominance check & candidates filtering:

B-tree Label Set:

Logical View:



Router keys are copies of the largest label in each subtree.